

IT'S TIME FOR SOME
RATE MATCHING!



Arun Raghavan

asymptotic

ME

- GStreamer, PipeWire, PulseAudio person
- Founder @ [asymptotic](#)
 - Help build cool things with these projects
 - Contribute back to upstream better
 - ???
 - Positive feedback loop

CLOCKS

- Devices derive time from a clock
- A repeating signal with a fixed frequency
- Generated by a quartz crystal



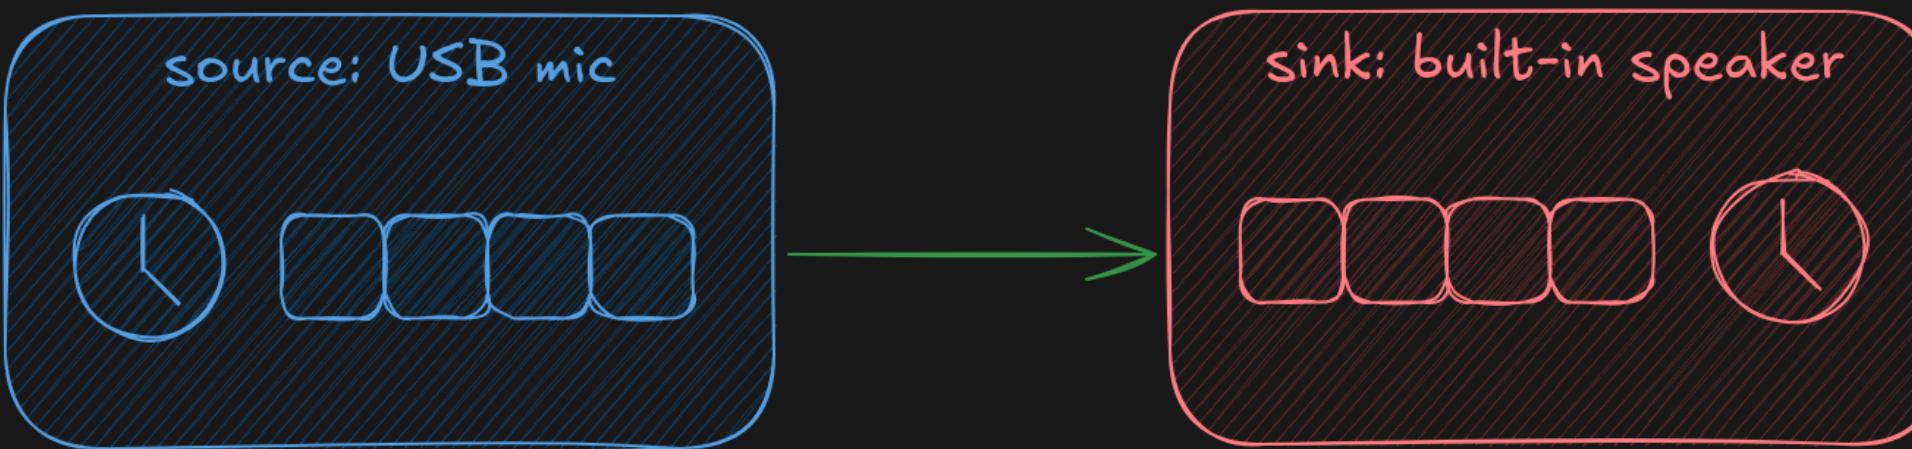
TIME IS AN ILLUSION

- Different devices will have different crystals
- With *slightly different* notions of time



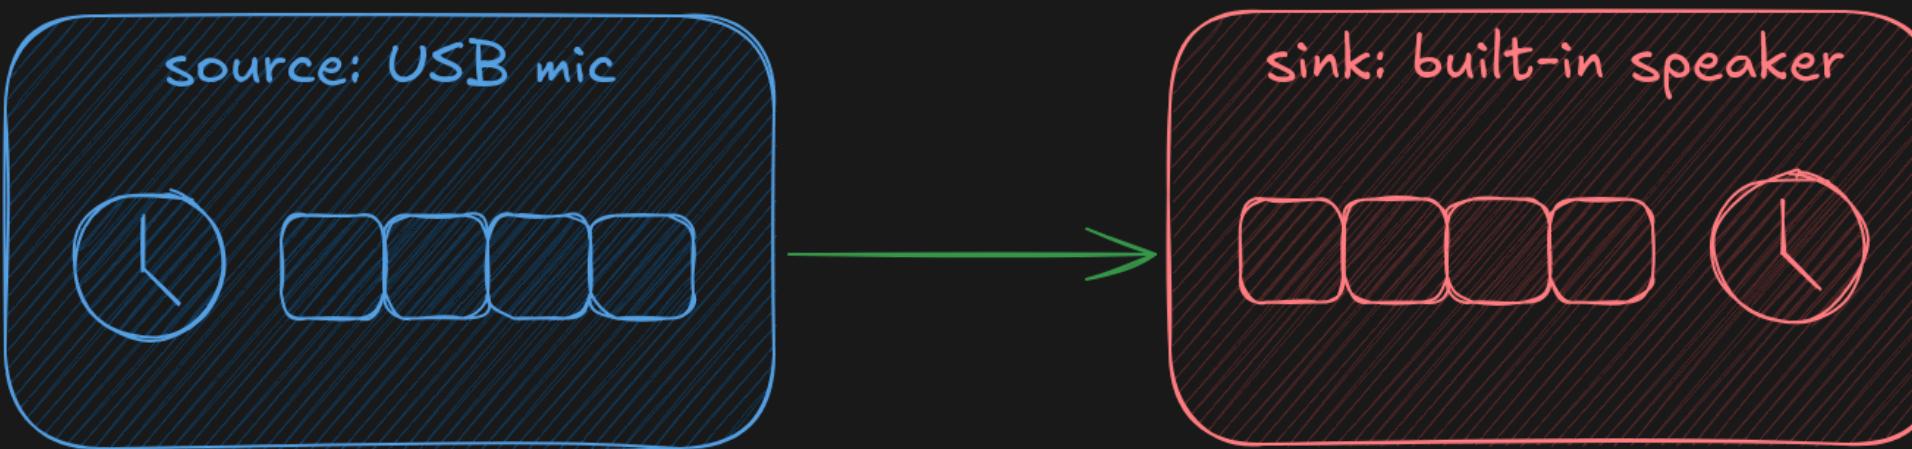
What does 1ms even mean?

THE RATE MATCHING PROBLEM



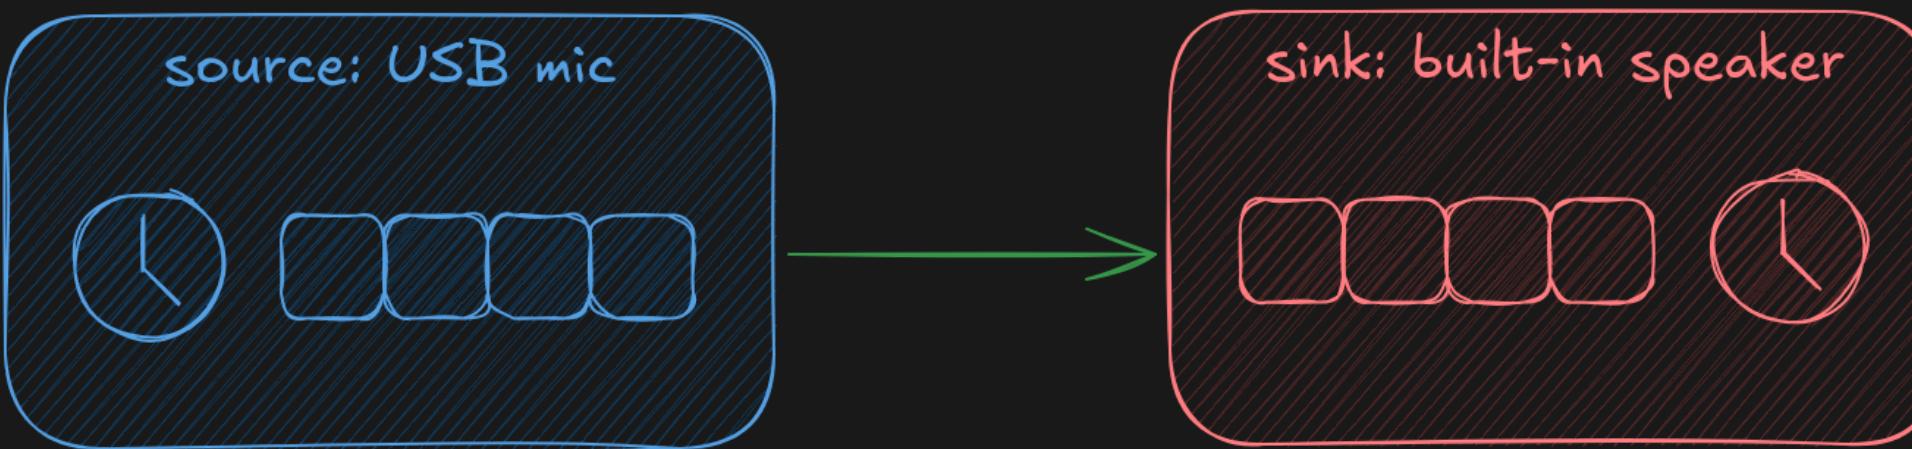
- Two devices with different clocks
- Producing data every “1ms”, say

THE RATE MATCHING PROBLEM



- What happens if the source is *slower* than sink
- Or if the source is *faster* than the sink?

THE RATE MATCHING PROBLEM



- Source slower than sink: *Underflow!*
- Sink slower than source: *Overflow!*

GSTREAMER

GSTREAMER: CLOCKS

- Clocks are represented via GstClock
- Report the current time
- Wait until a certain time

GSTREAMER: CLOCKS

- One for the pipeline
- Elements synchronise to this
- i.e. decide when to render data

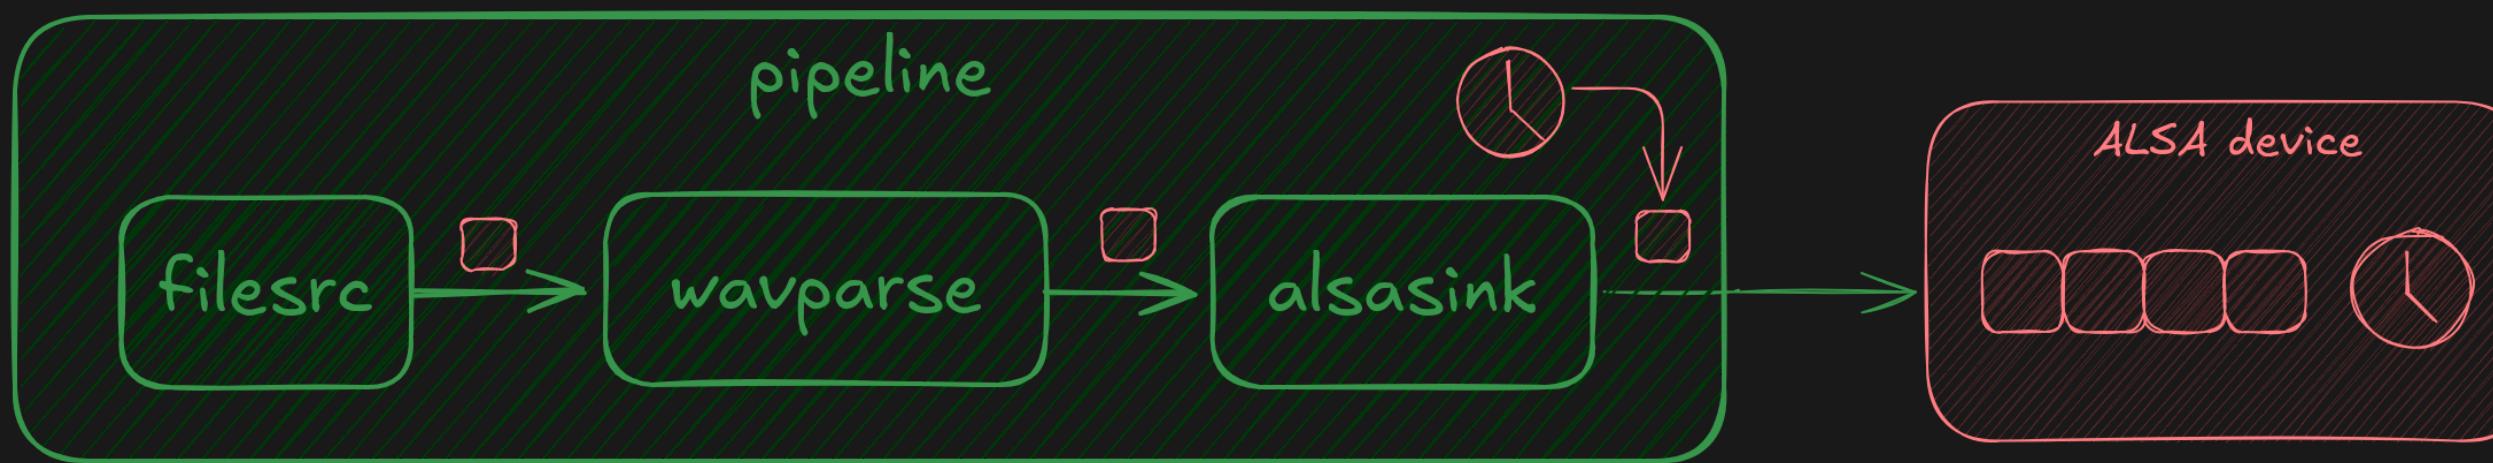
GSTREAMER: CLOCKS

- Elements can provide their own clocks
- System clock
- Audio clocks
- NTP, PTP, netclock

GSTREAMER: CLOCKS

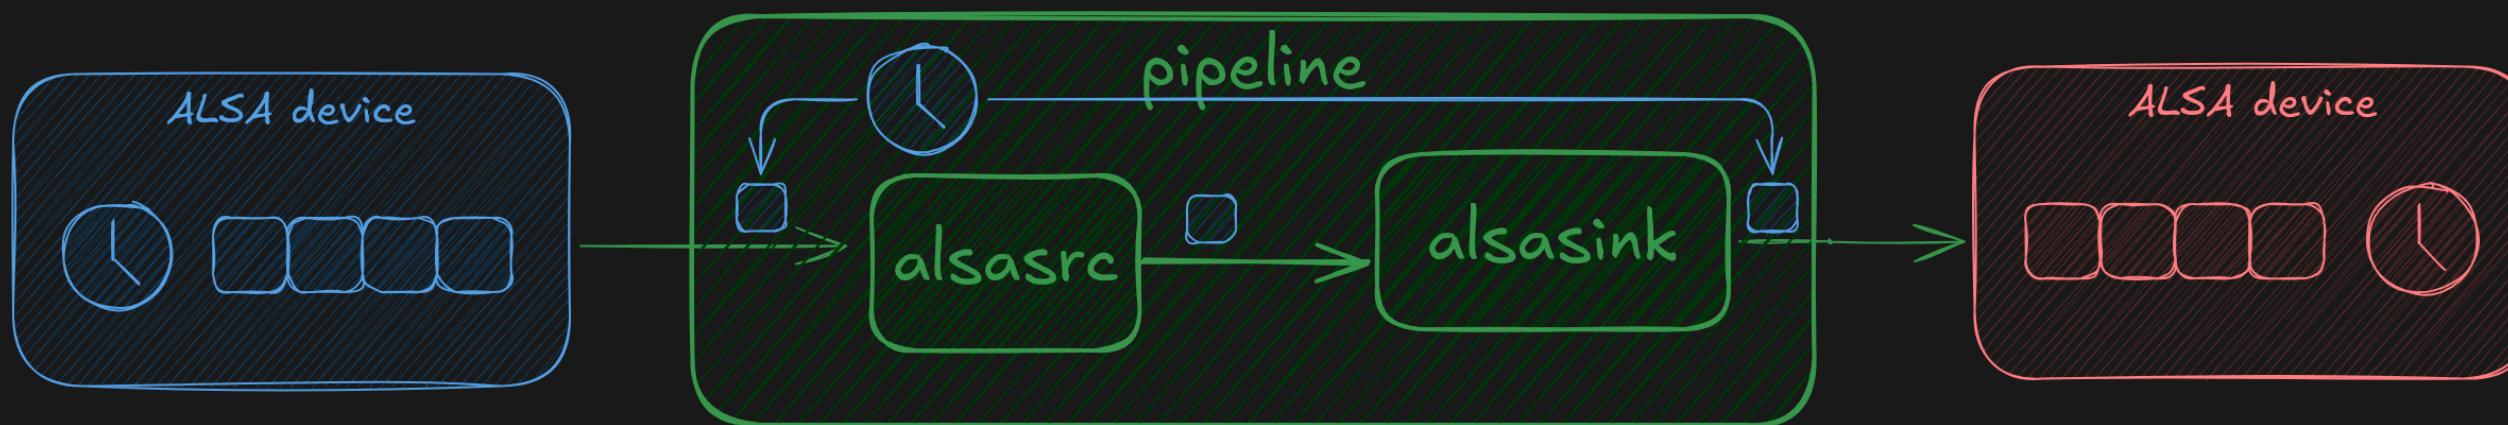
- Pair to system clock for continuous output
- Linear regression
 - $t_{audio}^1 \implies t_{sys}^1$
 - $t_{audio}^2 \implies t_{sys}^2$
 - $\therefore t_{sys}^{now} \implies t_{audio}^{now}$
- With some heuristic filtering

GSTREAMER: CLOCKS



A non-live pipeline

GSTREAMER: CLOCKS



A live pipeline

GSTREAMER: RATE MATCHING

- Only some elements sync to clock
- GstAudioBaseSink
- GstAudioBaseSrc
- GstAudioAggregator

GSTREAMER: RATE MATCHING

They may...

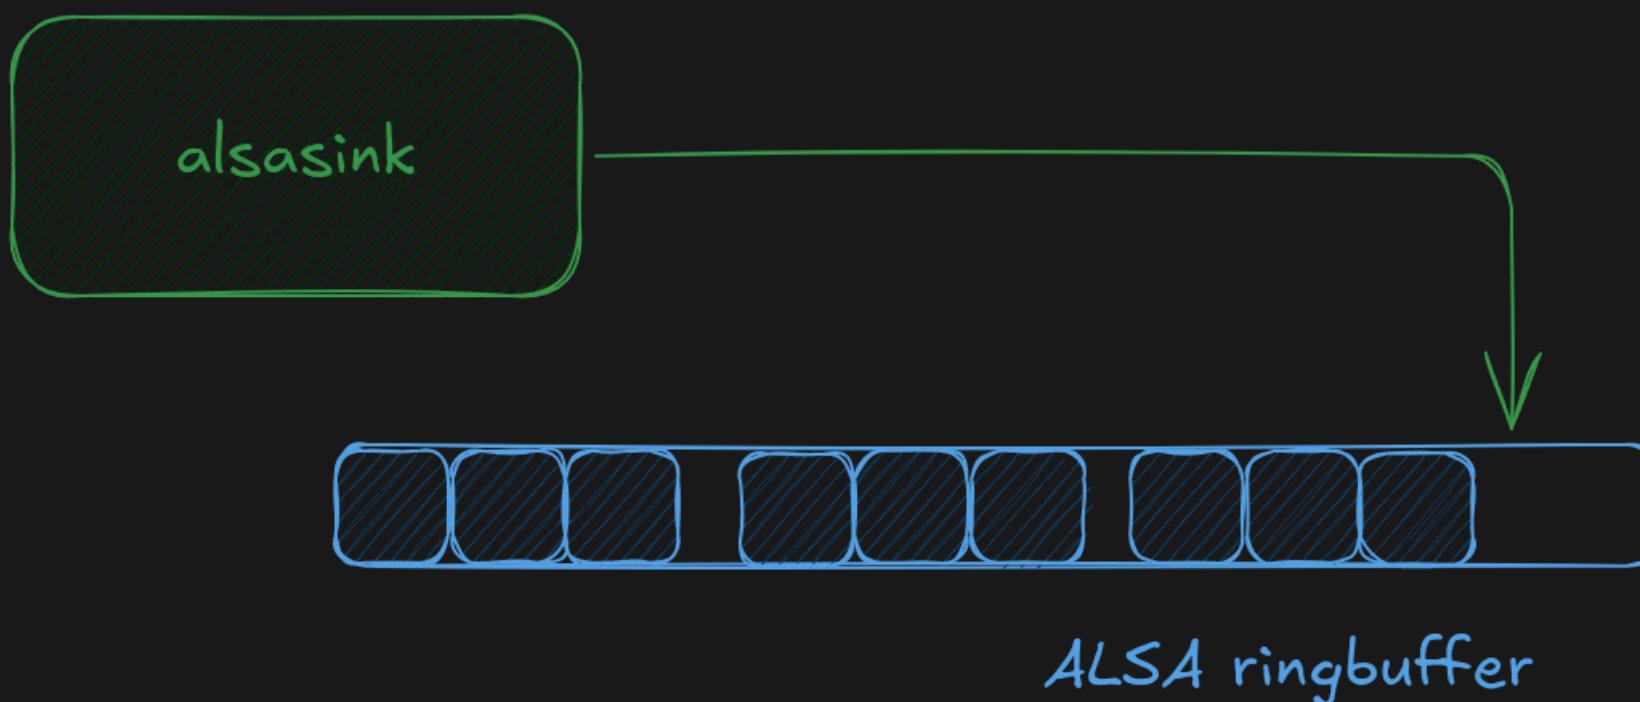
1. Skew
2. Resample
3. Custom

GSTREAMER: RATE MATCHING

- Skewing is most common
- “Spread” out the error
- Less audible glitches

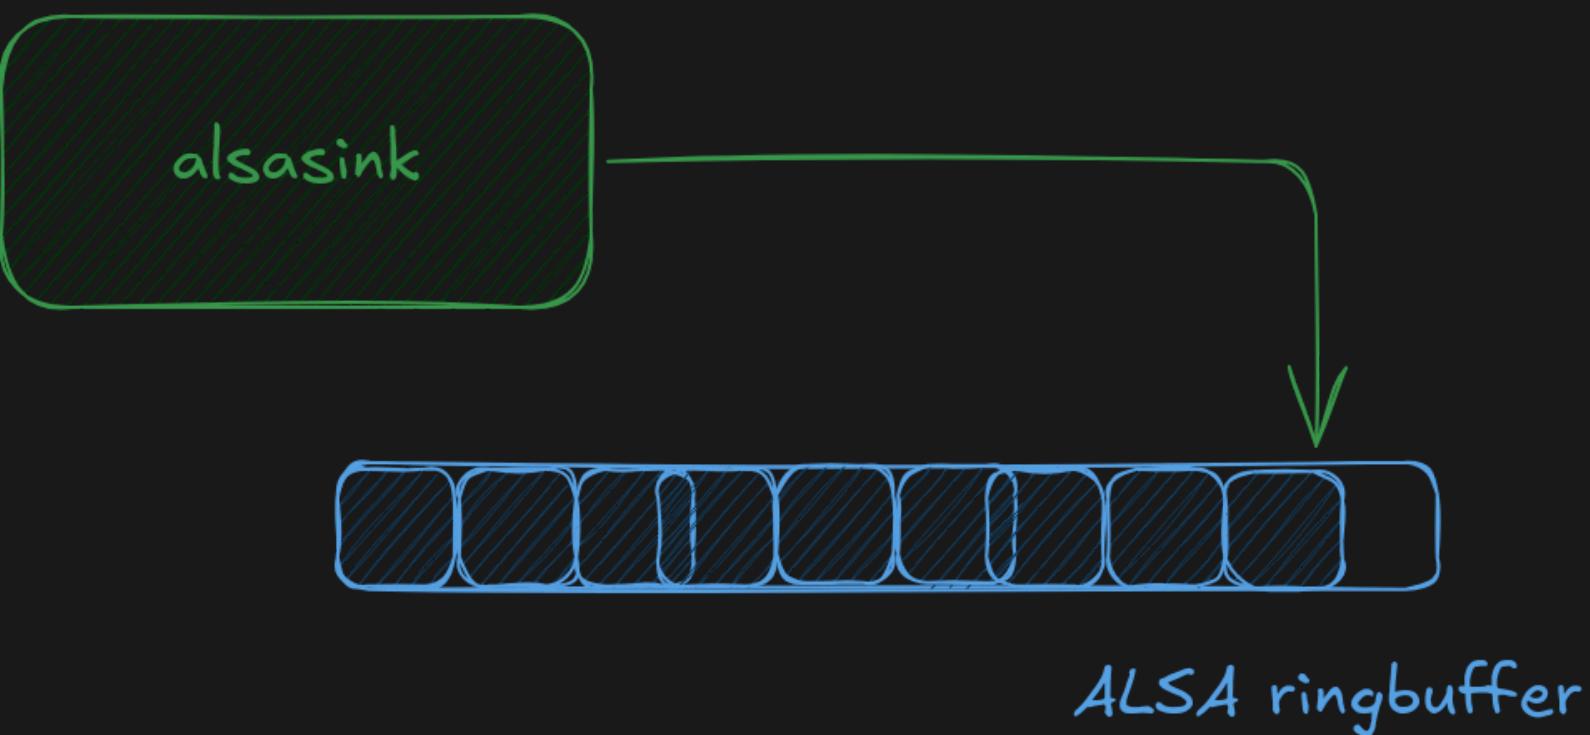
GSTREAMER: RATE MATCHING

- Sink faster than source
- Insert smaller gaps instead of underflowing



GSTREAMER: RATE MATCHING

- Sink slower than source
- Overwrite smaller chunks instead of overflowing



GSTREAMER: RATE MATCHING

- Resampling
- *Magical* mathematics

GSTREAMER: RATE MATCHING

- Produce the correct number of samples
- Implemented as linear interpolation
- Naïve and low quality
- Limitations of GstAudioRingBuffer

GSTREAMER: RATE MATCHING

- Custom
- Let the application do *something*
- Like adjust the audio clock itself
 - “Pull” a PLL

PIPEWIRE

PIPEWIRE: MODEL

- Also a graph
- process () cycle
- Quantum-sized buffers
- Much like JACK



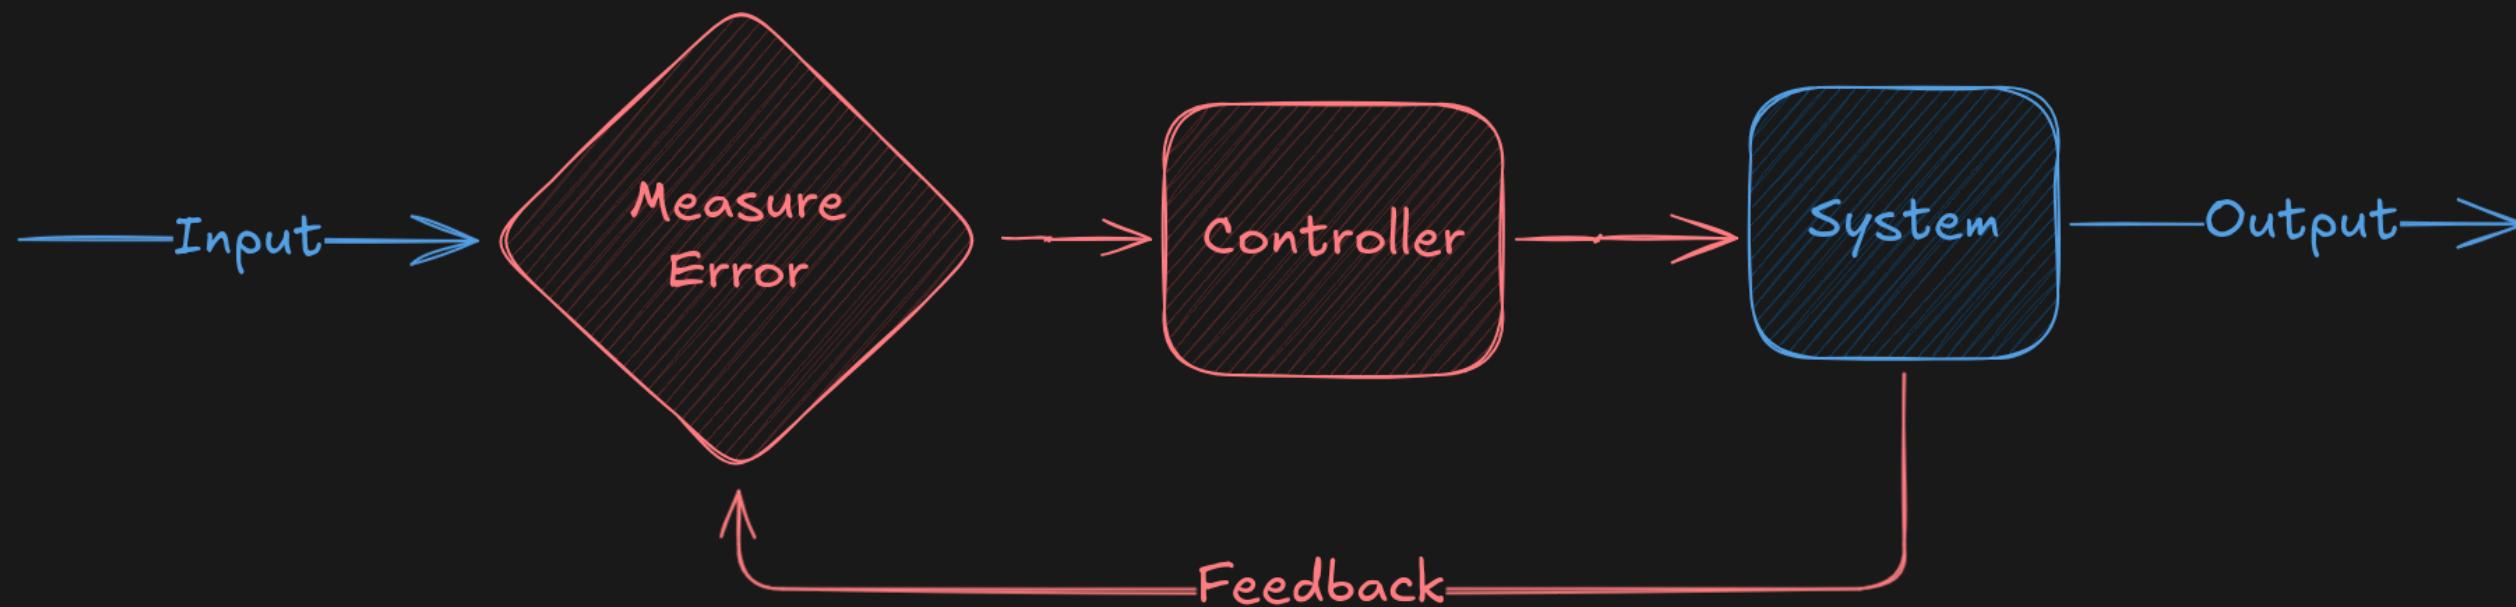
PIPEWIRE: DRIVERS

- One *driver* per graph
- All other nodes are *followers*
- Provides tick-per-quantum
- May or may not participate in data flow

PIPEWIRE: DRIVERS

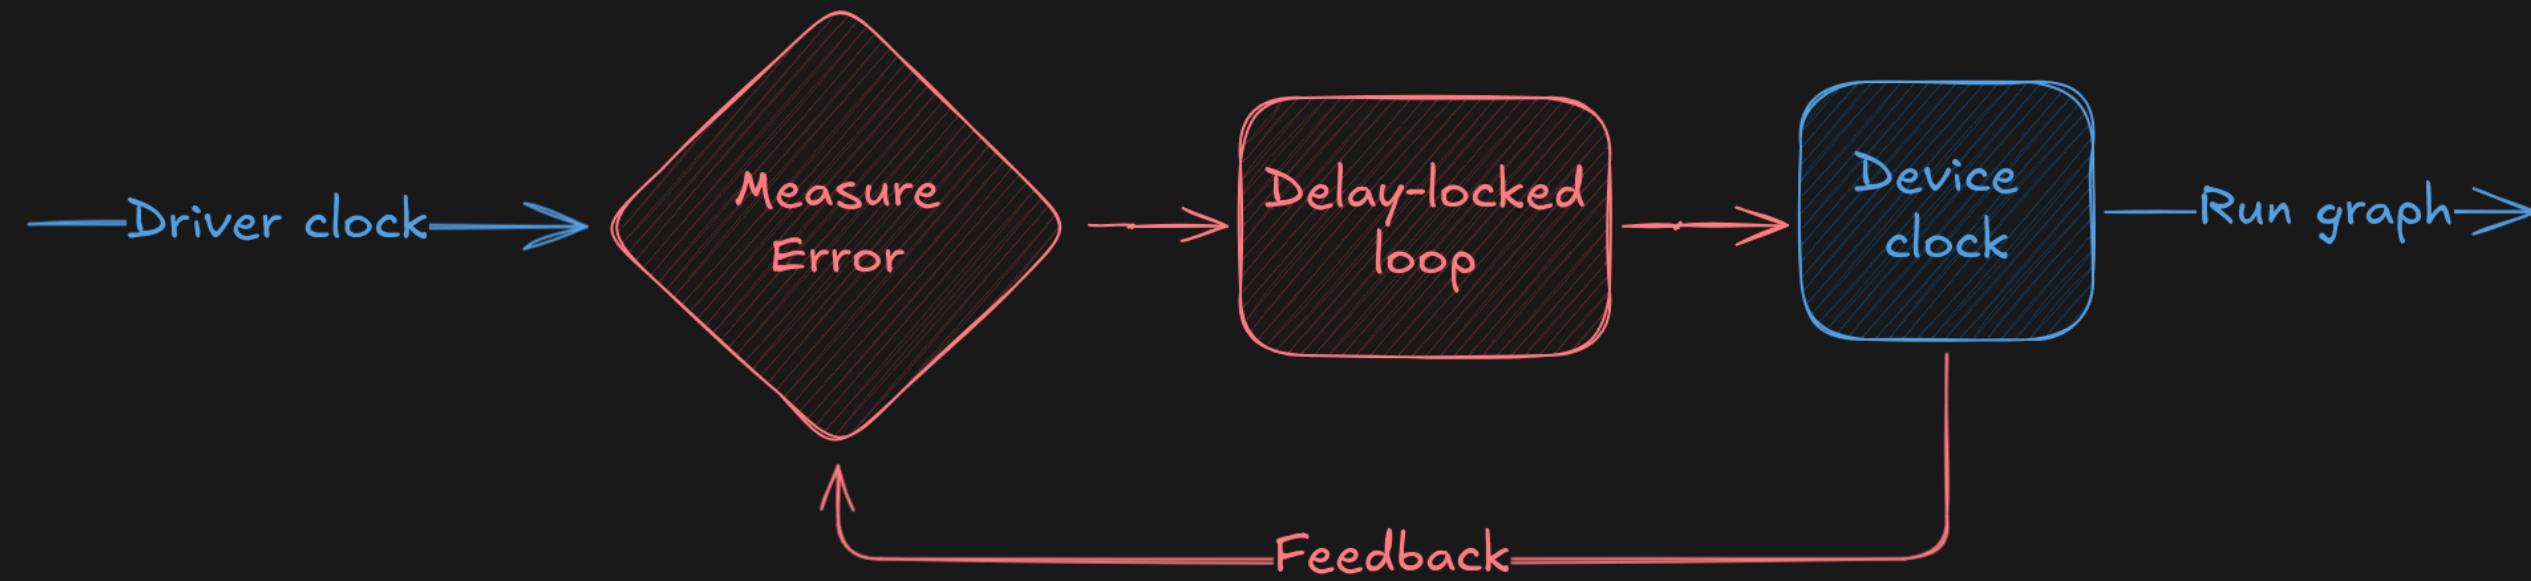
- System clock
- PTP
- Devices

DETOUR: CONTROL SYSTEMS



Closed loop control

PIPEWIRE: DELAY-LOCKED LOOPS

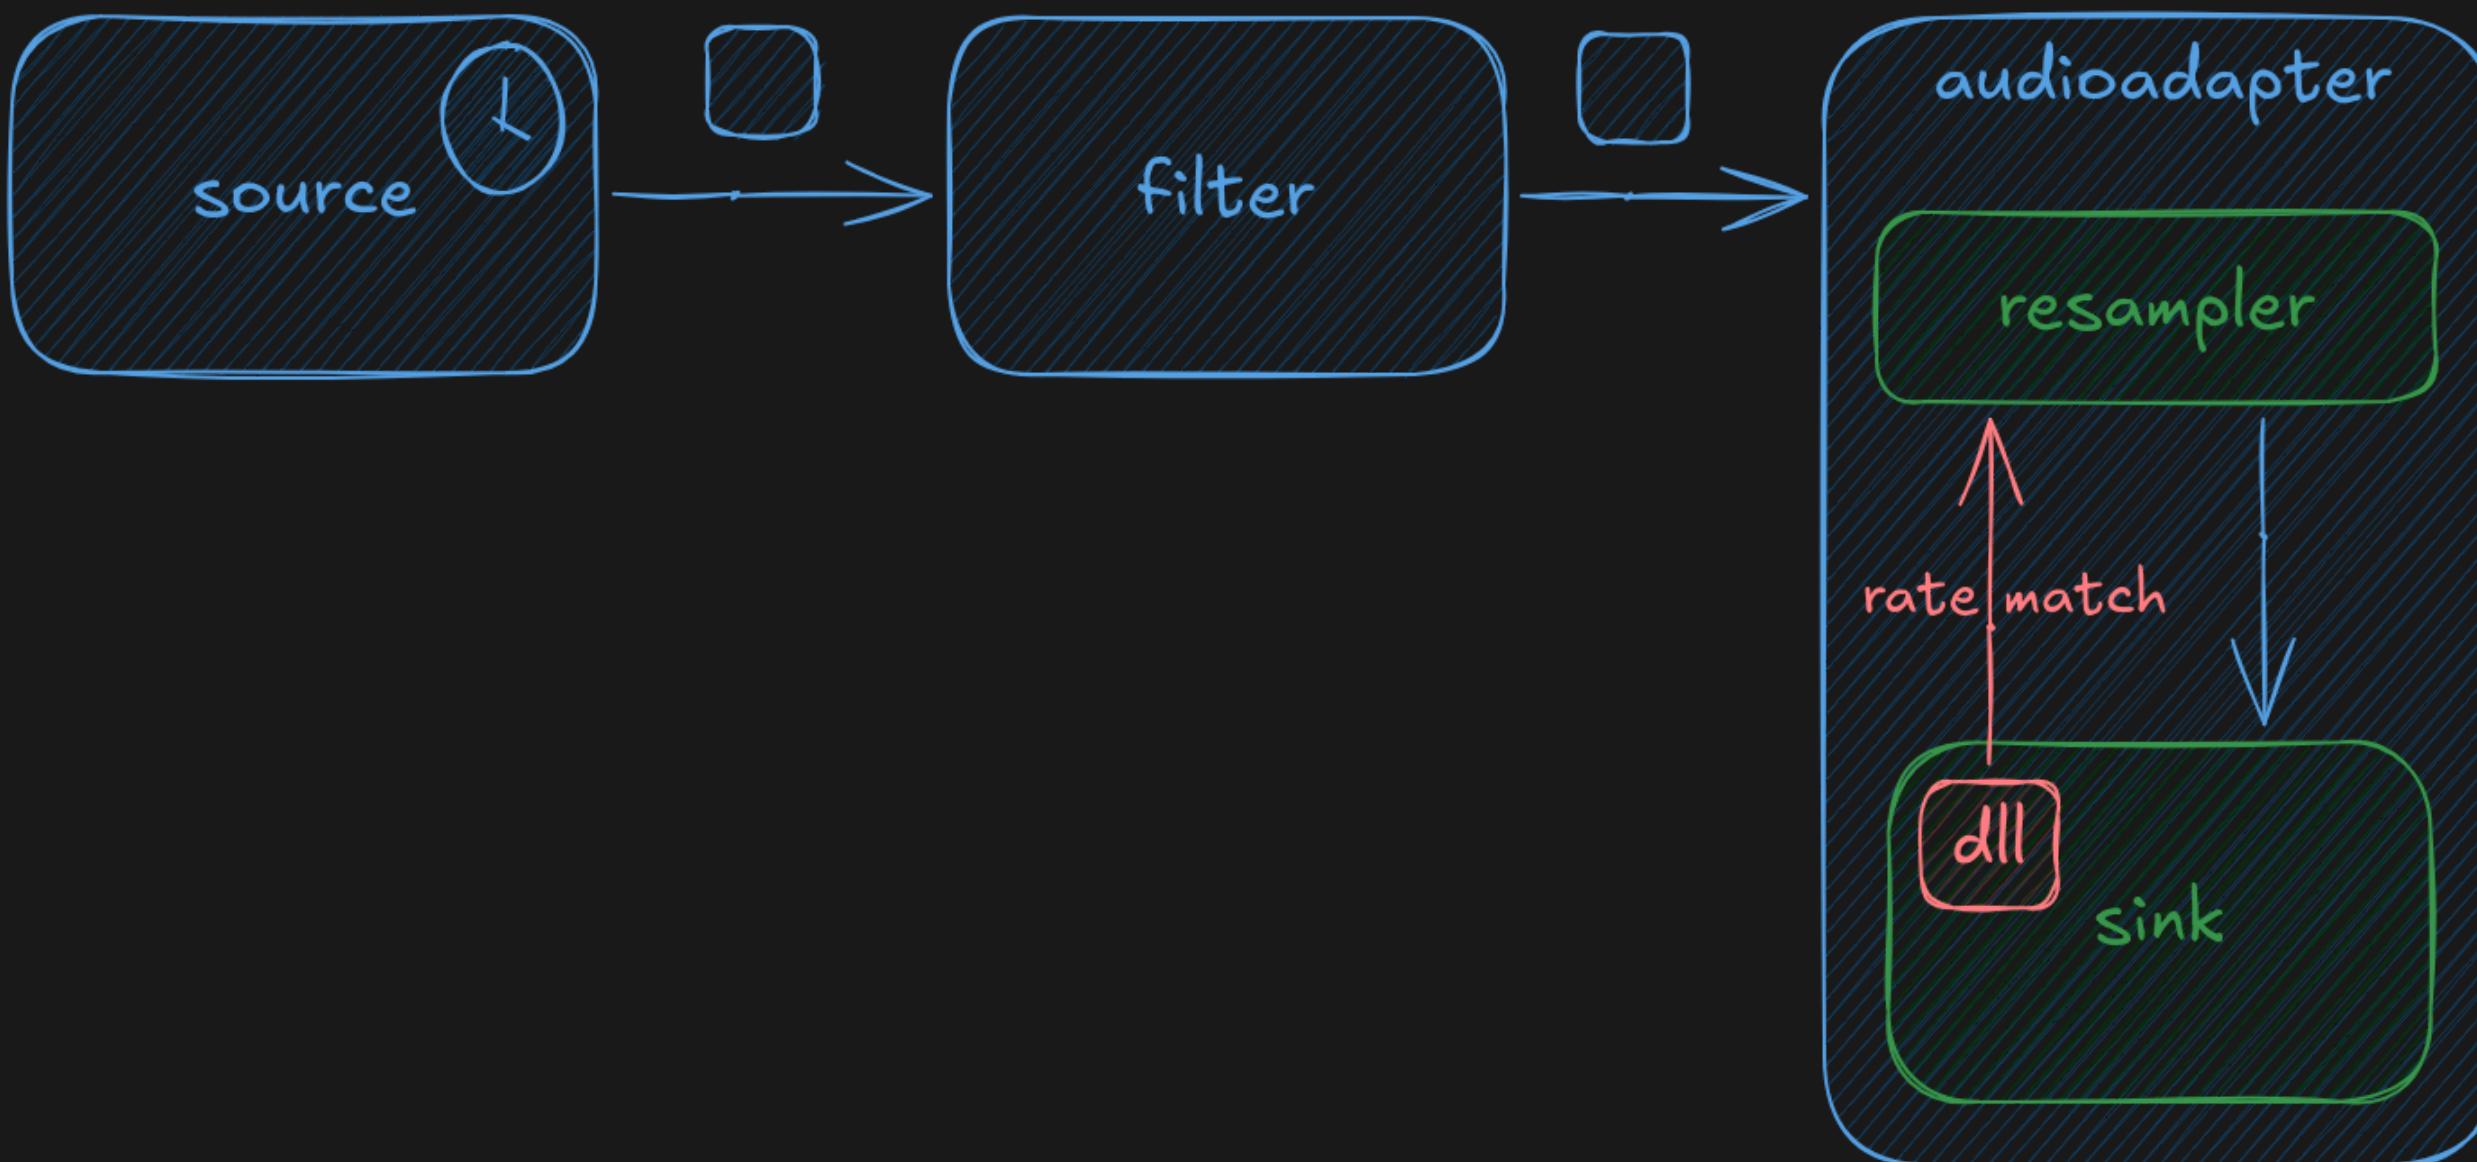


Rate matching as a control system

PIPEWIRE: DELAY-LOCKED LOOPS

- For each cycle, we have:
- nsec: Current cycle start time
- next_nsec: Expected next cycle start time
- *The relative rate of the driver and device clock*

PIPEWIRE: RATE MATCHING



PIPEWIRE: RATE MATCHING

- Can adapt to arbitrary clocks
- Can also adapt across graphs & mismatched quanta
- Resampler does fine-grained adjustment
- Converges quickly, hovers in a tight range
- Can use other mechanisms too

PIPEWIRE: RATE MATCHING REFERENCES

- Heavily inspired by JACK
- Bunch of work by Fons Adriaensen
 - Using a DLL to filter time
 - Controlling adaptive resampling

OBSERVATIONS

- System clock as a source for continuous time
- Open-ish (linear regression + heuristics) vs. closed control loops
 - Audio clocks work well
 - Network clock jitter can be problematic

OBSERVATIONS

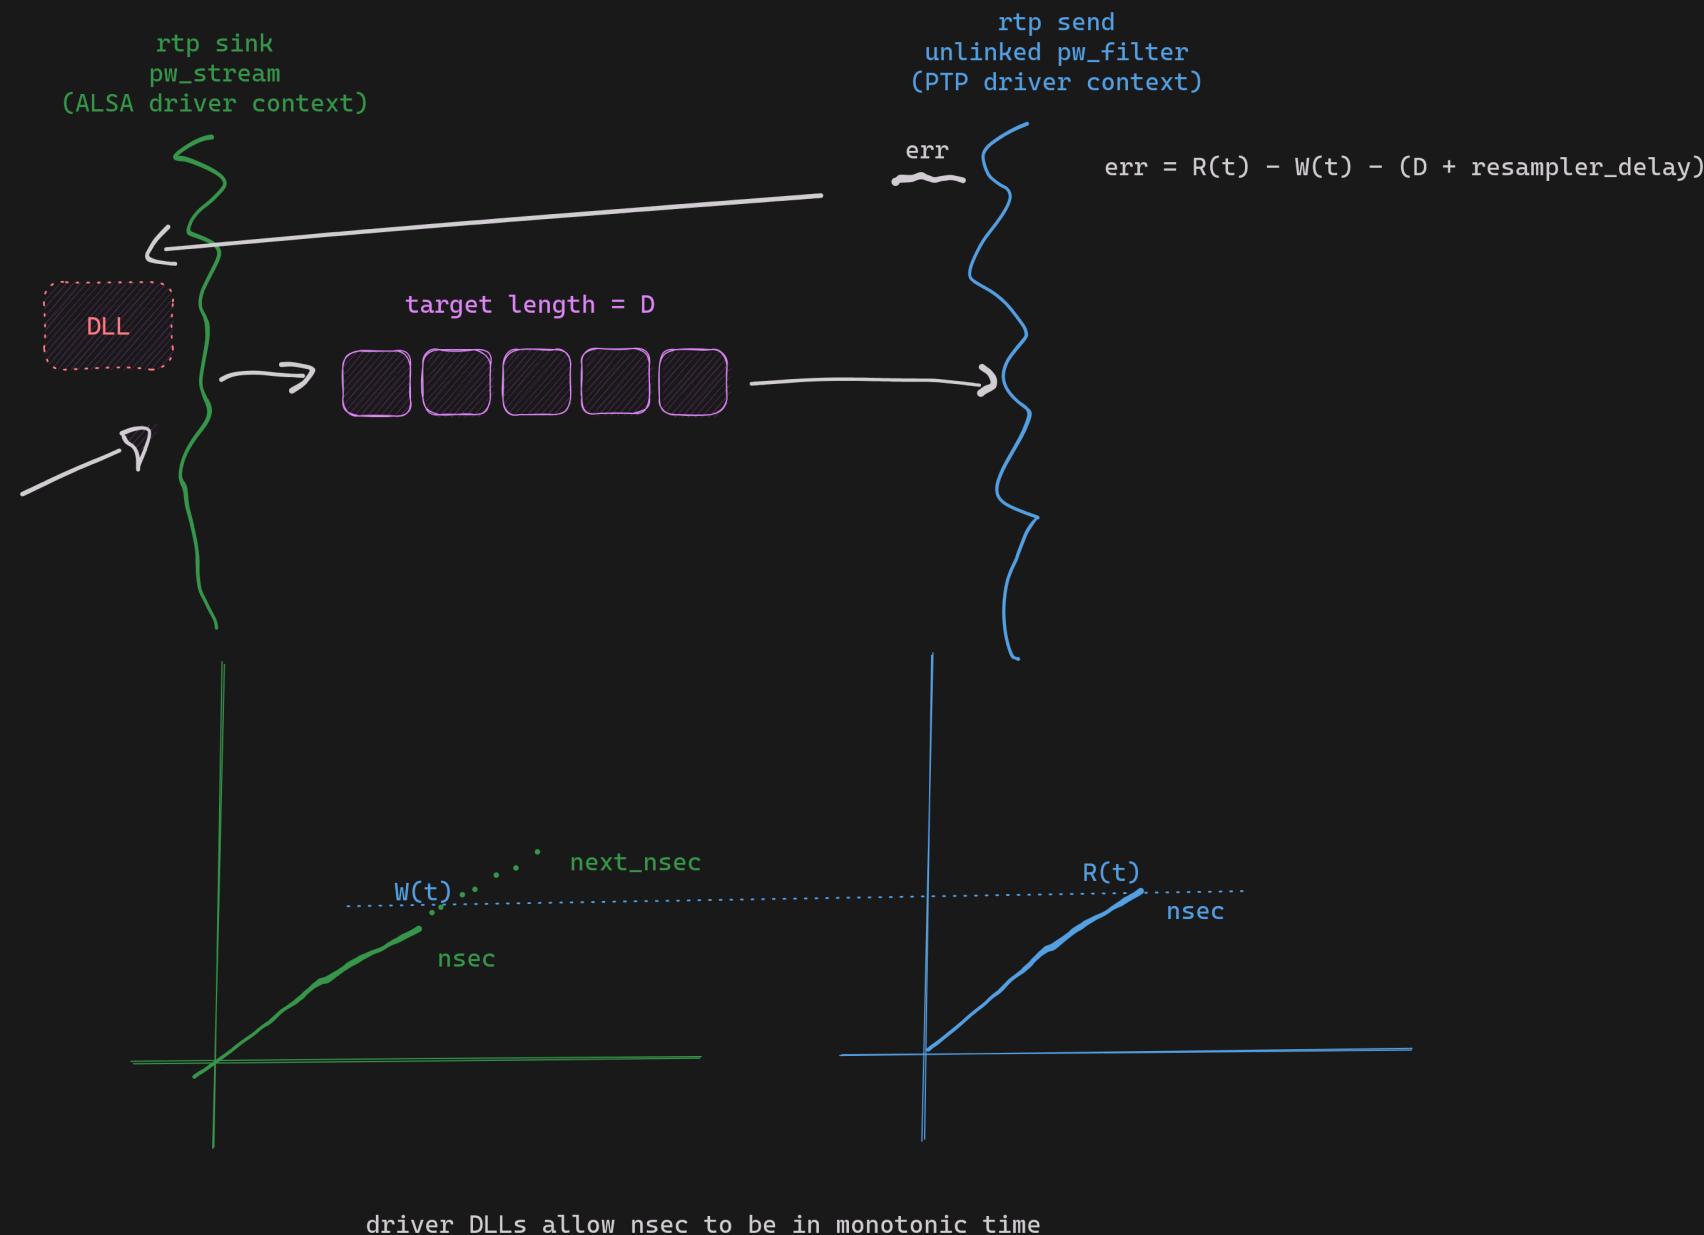
- You can skip GstAudioRingBuffer, like
 - Decklink elements
 - pipewireaudiosink
 - pulsedirectsink

QUESTIONS?



asymptotic

BONUS SLIDE



cross-graph rate matching